

▼ Stock Trading Using Deep Q-Learning

▼ Problem Statement

Prepare an agent by implementing Deep Q-Learning that can perform unsupervised trading in stock trade. The aim of this project is to train an agent that uses Q-learning and neural networks to predict the profit or loss by building a model and implementing it on a dataset that is available for evaluation.

The stock trading index environment provides the agent with a set of actions:

- Buy
- Sell
- Sit

This project has following sections:

- Import libraries
- Create a DQN agent
- Preprocess the data
- Train and build the model
- Evaluate the model and agent

Steps to perform

In the section **create a DQN agent**, create a class called agent where:

- Action size is defined as 3
- Experience replay memory to deque is 1000
- Empty list for stocks that has already been bought
- The agent must possess the following hyperparameters:
 - $\gamma = 0.95$
 - $\epsilon = 1.0$
 - $\epsilon_{\text{final}} = 0.01$
 - $\epsilon_{\text{decay}} = 0.995$

Note: It is advised to compare the results using different values in hyperparameters.

- Neural network has 3 hidden layers
- Action and experience replay are defined

▼ Solution

▼ Import the libraries

```
from google.colab import drive

drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.moun

```
import keras
from keras.models import Sequential
from keras.models import load_model
from keras.layers import Dense
from keras.optimizers import adam_v2
import numpy as np
import random
from collections import deque
```

▼ Create a DQN agent

Use the instruction below to prepare an agent

```
# Action space include 3 actions: Buy, Sell, and Sit
#Setting up the experience replay memory to deque with 1000 elements inside it
#Empty list with inventory is created that contains the stocks that were already bought
#Setting up gamma to 0.95, that helps to maximize the current reward over the long-term
#Epsilon parameter determines whether to use a random action or to use the model for the acti
#In the beginning random actions are encouraged, hence epsilon is set up to 1.0 when the mode
#And over time the epsilon is reduced to 0.01 in order to decrease the random actions and use
#We're then set the speed of decreasing epsililon in the epsilon_decay parameter

#Defining our neural network:
#Define the neural network function called _model and it just takes the keyword self
#Define the model with Sequential()
#Define states i.e. the previous n days and stock prices of the days
```

```
#Defining 3 hidden layers in this network
```

```
#Changing the activation function to relu because mean-squared error is used for the loss
```

```
class Agent:
```

```
    def __init__(self, state_size, is_eval=False, model_name=""):
        self.state_size = state_size # normalized previous days
        self.action_size = 3 # sit, buy, sell
        self.memory = deque(maxlen=1000)
        self.inventory = []
        self.model_name = model_name
        self.is_eval = is_eval
        self.gamma = 0.95
        self.epsilon = 1.0
        self.epsilon_min = 0.01
        self.epsilon_decay = 0.995
        self.model = load_model(model_name) if is_eval else self._model()
```

```
    def _model(self):
        model = Sequential()
        model.add(Dense(units=64, input_dim=self.state_size, activation="relu"))
        model.add(Dense(units=32, activation="relu"))
        model.add(Dense(units=8, activation="relu"))
        model.add(Dense(self.action_size, activation="linear"))
        model.compile(loss="mse", optimizer=adam_v2.Adam(lr=0.001))
        return model
```

```
    def act(self, state):
        if not self.is_eval and random.random() <= self.epsilon:
            return random.randrange(self.action_size)
        options = self.model.predict(state)
        return np.argmax(options[0])
```

```
    def expReplay(self, batch_size):
        mini_batch = []
        l = len(self.memory)
        for i in range(l - batch_size + 1, l):
            mini_batch.append(self.memory[i])
        for state, action, reward, next_state, done in mini_batch:
            target = reward
            if not done:
                target = reward + self.gamma * np.amax(self.model.predict(next_state)[0])
            target_f = self.model.predict(state)
            target_f[0][action] = target
            self.model.fit(state, target_f, epochs=1, verbose=0)
        if self.epsilon > self.epsilon_min:
            self.epsilon *= self.epsilon_decay
```

▼ Preprocess the stock market data

```
import math

# prints formatted price
def formatPrice(n):
    return ("-$" if n < 0 else "$") + "{0:.2f}".format(abs(n))

# returns the vector containing stock data from a fixed file
def getStockDataVec(key):
    vec = []
    lines = open("'" + key + ".csv", "r").read().splitlines()

    for line in lines[1:]:
        vec.append(float(line.split(",")[4]))

    return vec

# returns the sigmoid
def sigmoid(x):
    return 1 / (1 + math.exp(-x))

# returns an n-day state representation ending at time t
def getState(data, t, n):
    d = t - n + 1
    block = data[d:t + 1] if d >= 0 else -d * [data[0]] + data[0:t + 1] # pad with t0
    res = []
    for i in range(n - 1):
        res.append(sigmoid(block[i + 1] - block[i]))

    return np.array([res])
```

▼ Train and build the model

```
'''import sys

if len(sys.argv) != 4:
    print ("Usage: python train.py [stock] [window] [episodes]")
    exit()

stock_name = input("Enter stock_name, window_size, Episode_count")
#Fill the given information when prompted:
#Enter stock_name = GSPC_Training_Dataset
#window_size = 10
```

```

#Episode_count = 100 or it can be 10 or 20 or 30 and so on.

window_size = input()
episode_count = input()
stock_name = str(stock_name)
window_size = int(window_size)
episode_count = int(episode_count)

agent = Agent(window_size)
data = getStockDataVec(stock_name)
l = len(data) - 1
batch_size = 32

for e in range(episode_count + 1):
    print ("Episode " + str(e) + "/" + str(episode_count))
    state = getState(data, 0, window_size + 1)

    total_profit = 0
    agent.inventory = []

    for t in range(l):
        action = agent.act(state)

        # sit
        next_state = getState(data, t + 1, window_size + 1)
        reward = 0

        if action == 1: # buy
            agent.inventory.append(data[t])
            print ("Buy: " + formatPrice(data[t]))

        elif action == 2 and len(agent.inventory) > 0: # sell
            bought_price = agent.inventory.pop(0)
            reward = max(data[t] - bought_price, 0)
            total_profit += data[t] - bought_price
            print ("Sell: " + formatPrice(data[t]) + " | Profit: " + formatPrice(data[t] - bo

        done = True if t == l - 1 else False
        agent.memory.append((state, action, reward, next_state, done))
        state = next_state

    if done:
        print ("-----")
        print ("Total Profit: " + formatPrice(total_profit))

    if len(agent.memory) > batch_size:
        agent.expReplay(batch_size)

    #if e % 10 == 0:
        agent.model.save("model_ep" + str(e))'''

```

```

import sys

if len(sys.argv) != 4:
    print ("Usage: python train.py [stock] [window] [episodes]")
    exit()

stock_name = input("Enter stock_name, window_size, Episode_count")
#Fill the given information when prompted:
#Enter stock_name = GSPC_Training_Dataset
#window_size = 10
#Episode_count = 100 or it can be 10 or 20 or 30 and so on.

window_size = input()
episode_count = input()
stock_name = str(stock_name)
window_size = int(window_size)
episode_count = int(episode_count)

agent = Agent(window_size)
data = getStockDataVec(stock_name)
l = 10
batch_size = 32

for e in range(episode_count + 1):
    print ("Episode " + str(e) + "/" + str(episode_count))
    state = getState(data, 0, window_size + 1)

    total_profit = 0
    agent.inventory = []

    for t in range(l):
        action = agent.act(state)

        # sit
        next_state = getState(data, t + 1, window_size + 1)
        reward = 0

        if action == 1: # buy
            agent.inventory.append(data[t])
            print ("Buy: " + formatPrice(data[t]))

        elif action == 2 and len(agent.inventory) > 0: # sell
            bought_price = agent.inventory.pop(0)
            reward = max(data[t] - bought_price, 0)
            total_profit += data[t] - bought_price
            print ("Sell: " + formatPrice(data[t]) + " | Profit: " + formatPrice(data[t] - bo

```

```

done = True if t == 1 - 1 else False
agent.memory.append((state, action, reward, next_state, done))
state = next_state

if done:
    print ("-----")
    print ("Total Profit: " + formatPrice(total_profit))

if len(agent.memory) > batch_size:
    agent.expReplay(batch_size)

# if e % 10 == 0:
    agent.model.save("model_ep" + str(e))

```

Episode 8/10

INFO:tensorflow:Assets written to: model_ep8/assets

INFO:tensorflow:Assets written to: model_ep8/assets

Buy: \$1333.34

INFO:tensorflow:Assets written to: model_ep8/assets

Buy: \$1298.35

INFO:tensorflow:Assets written to: model_ep8/assets

INFO:tensorflow:Assets written to: model_ep8/assets

Sell: \$1300.80 | Profit: -\$32.54

INFO:tensorflow:Assets written to: model_ep8/assets

Sell: \$1313.27 | Profit: \$14.92

INFO:tensorflow:Assets written to: model_ep8/assets

Buy: \$1326.82

INFO:tensorflow:Assets written to: model_ep8/assets

Buy: \$1318.55

INFO:tensorflow:Assets written to: model_ep8/assets

Sell: \$1326.65 | Profit: -\$0.17

Total Profit: -\$17.79

INFO:tensorflow:Assets written to: model_ep8/assets

Episode 9/10

INFO:tensorflow:Assets written to: model_ep9/assets

INFO:tensorflow:Assets written to: model_ep9/assets

Buy: \$1333.34

INFO:tensorflow:Assets written to: model_ep9/assets

INFO:tensorflow:Assets written to: model_ep9/assets

INFO:tensorflow:Assets written to: model_ep9/assets

Buy: \$1300.80

INFO:tensorflow:Assets written to: model_ep9/assets

Buy: \$1313.27

INFO:tensorflow:Assets written to: model_ep9/assets

Sell: \$1326.82 | Profit: -\$6.52

INFO:tensorflow:Assets written to: model_ep9/assets

Buy: \$1318.55

INFO:tensorflow:Assets written to: model_ep9/assets

Buy: \$1326.65

Total Profit: -\$6.52

INFO:tensorflow:Assets written to: model_ep9/assets

Episode 10/10

Buy: \$1300.80

```

buy: 1400.21
INFO:tensorflow:Assets written to: model_ep10/assets
INFO:tensorflow:Assets written to: model_ep10/assets
Sell: $1333.34 | Profit: $50.07
INFO:tensorflow:Assets written to: model_ep10/assets
Buy: $1298.35
INFO:tensorflow:Assets written to: model_ep10/assets
Sell: $1295.86 | Profit: -$2.49
INFO:tensorflow:Assets written to: model_ep10/assets
INFO:tensorflow:Assets written to: model_ep10/assets
INFO:tensorflow:Assets written to: model_ep10/assets
INFO:tensorflow:Assets written to: model_ep10/assets
Buy: $1318.55
INFO:tensorflow:Assets written to: model_ep10/assets
-----
Total Profit: $47.58
INFO:tensorflow:Assets written to: model_ep10/assets

```

▼ Evaluate the model and agent

```

'''import sys
from keras.models import load_model

if len(sys.argv) != 3:
    print ("Usage: python evaluate.py [stock] [model]")
    exit()

stock_name = input("Enter Stock_name, Model_name")
model_name = input()
#Note:
#Fill the given information when prompted:
#Enter stock_name = GSPC_Evaluation_Dataset
#Model_name = respective model name

model = load_model("'" + model_name)
window_size = model.layers[0].input.shape.as_list()[1]

agent = Agent(window_size, True, model_name)
data = getStockDataVec(stock_name)
l = len(data) - 1
batch_size = 32

state = getState(data, 0, window_size + 1)
total_profit = 0
agent.inventory = []

for t in range(l):
    action = agent.act(state)

```



```

# sit
next_state = getState(data, t + 1, window_size + 1)
reward = 0

if action == 1: # buy
    agent.inventory.append(data[t])
    print ("Buy: " + formatPrice(data[t]))

elif action == 2 and len(agent.inventory) > 0: # sell
    bought_price = agent.inventory.pop(0)
    reward = max(data[t] - bought_price, 0)
    total_profit += data[t] - bought_price
    print ("Sell: " + formatPrice(data[t]) + " | Profit: " + formatPrice(data[t] - bought_pri

done = True if t == l - 1 else False
agent.memory.append((state, action, reward, next_state, done))
state = next_state

if done:
    print ("-----")
    print (stock_name + " Total Profit: " + formatPrice(total_profit))'''

import sys
from keras.models import load_model

if len(sys.argv) != 3:
    print ("Usage: python evaluate.py [stock] [model]")
    exit()

stock_name = input("Enter Stock_name, Model_name")
model_name = input()
#Note:
#Fill the given information when prompted:
#Enter stock_name = GSPC_Evaluation_Dataset
#Model_name = respective model name

model = load_model("'" + model_name)
window_size = model.layers[0].input.shape.as_list()[1]

agent = Agent(window_size, True, model_name)
data = getStockDataVec(stock_name)
l = 10
batch_size = 32

state = getState(data, 0, window_size + 1)
total_profit = 0

```

```

agent.inventory = []

for t in range(1):
    action = agent.act(state)

    # sit
    next_state = getState(data, t + 1, window_size + 1)
    reward = 0

    if action == 1: # buy
        agent.inventory.append(data[t])
        print ("Buy: " + formatPrice(data[t]))

    elif action == 2 and len(agent.inventory) > 0: # sell
        bought_price = agent.inventory.pop(0)
        reward = max(data[t] - bought_price, 0)
        total_profit += data[t] - bought_price
        print ("Sell: " + formatPrice(data[t]) + " | Profit: " + formatPrice(data[t] - bought_price))

    done = True if t == 1 - 1 else False
    agent.memory.append((state, action, reward, next_state, done))
    state = next_state

    if done:
        print ("-----")
        print (stock_name + " Total Profit: " + formatPrice(total_profit))

```

Enter Stock_name, Model_name/content/drive/MyDrive/Colab Notebooks/GSPC_Evaluation_Dataset
/content/model_ep10

/content/drive/MyDrive/Colab Notebooks/GSPC_Evaluation_Dataset Total Profit: \$0.00



Note: Run the training section for considerable episodes so that while evaluating the model it can generate significant profit.

✓ 32s completed at 10:12 AM

● ×